What would have proved us wrong

2 October onward. We wrote failure conditions, but our validation remained a simple simulation. The strongest corrections came from evidence that challenged the framing and the numbers.

What actually broke in our reasoning

Our broad claim that texting prevents needs from being shared did not survive the full assistant account. Trusted telephone coordination could already be efficient, while some information was deliberately withheld. A new interface could add effort or reduce trust instead of helping.

Our “no standard SOP” claim also exceeded the evidence. MO’s supplied material already specifies a preference-to-action framework and minimum guidance. We had to narrow the contribution to a proposed assistant expression flow and a more explicit situational-care coordination layer.

The numerical outcome story broke at the evidence boundary. Survey nonresponse, a small review sample and assumed retention uplift could not establish hotel performance. The case therefore reports these as withdrawn claims or untested targets.

What our simulation did not establish

We did not conduct real post-design user tests or observe live service delivery. We have no reliable task-by-task measurements of time saved, misunderstandings, staff effort or fulfilment failures. Functional interaction checks and a team walkthrough cannot establish desirability or operational feasibility.

The conditions under which we would stop

User side	We would stop or redesign if target assistants preferred the existing channel, refused the interactive flow, or felt it exposed information they did not want recorded.
Business side	We would stop or reduce scope if the flow failed to produce usable, confirmed requests, or raised total staff effort beyond an agreed capacity and cost limit.
The join	We would reject the service claim if more information was collected but arrangements were not delivered as agreed, or the resulting care felt unwanted or failed to improve recognition.

How we would test those conditions next

We would ask real assistants to complete the same request through the existing channel and prototype, then ask property staff to simulate the complete handoff. We would record errors, omitted needs, clarification, total effort and delivery decisions. The hotel and research team must agree failure thresholds before the session.

Our original Class 10 sheet also tied success to return rate. We now treat loyalty as a later hypothesis: an unchanged return rate alone would not identify which part of the service failed.