The brief and our reframing

4 September–12 October 2026. We explored how pre-arrival care could help guests feel recognised while keeping the hotel's promises deliverable.

The brief as given

We began with a tension in luxury hospitality: asking about a preference creates an expectation. More questions can multiply combinations of room preparation, timing and service requests, while staff capacity, room changes and cost to serve remain constrained. The brief described an earlier hyper-personalisation pilot that had been stopped because of operational complexity.

We were asked to build engagement and manage expectations after booking and before arrival, without over-promising. The commercial intent was brand equity, loyalty and spending on property. This was not a booking-conversion brief.

The brief as we reframed it

We are solving for pre-arrival care that helps guests feel recognised and can be delivered reliably at manageable cost and complexity, in service of Mandarin Oriental's goal of strengthening brand equity and guests' willingness to pay.

We chose an urban MO context and a high-net-worth guest segment, often represented by a human assistant. The "top 5%" boundary was our design focus, not a verified segment size. We connected what the assistant knows before arrival with what a Dream Maker may notice in person, then asked how both could become an owned, feasible service response.

How the work unfolded

Period	What we worked on
4–7 Sep	Interrogate the brief; separate facts, assumptions and questions; hear Kate's account.
11–19 Sep	Build evidence through fieldwork, MO conversations and assistant interviews; challenge the initial framing.
18–25 Sep	Develop insights, compare concepts, map stakeholders and work through the service journey.
28 Sep–2 Oct	Connect user value to business measures; write criteria and failure conditions.
By 5–12 Oct	Iterate the room prototype and Dream Maker workflow; assemble the final proposition and playback.

The timeline follows dated course and project records. Some workshop artefacts are undated; their exact production day is not inferred.

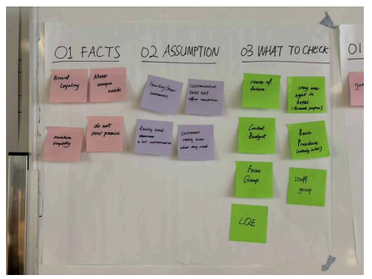

Fig. 1. Our early Facts / Assumptions / Checks wall. The case study follows the work behind the conclusion.

Evidence, access and the voices we missed

We combined proxy expertise, small field samples and assistant accounts. These inputs revealed a service opportunity, but they did not establish a universal communication problem.

Who we heard from

We first heard from Kate, our MO proxy stakeholder, about the promise created by preference questions and the operational burden of fulfilment. Further MO conversations helped us understand agent-mediated booking and the staffing constraints around high-touch service.

Our final research record reports five short conversations with guests matching our chosen profile at Hong Kong MO on 12 September; four said an assistant booked their stay. It also summarises three personal-agent interviews. These were small, accessible samples rather than representative research.

The Shenzhen MO visit introduced us to the Dream Maker role: a person who notices contextual signals and turns them into acts of care. This observation expanded the work beyond collecting information before arrival.

A useful counterexample, 19 September

One executive assistant described familiar telephone coordination as effective. She was wary of putting private needs into another platform and wanted an accountable manager to acknowledge requests. She was more positive about optional menus, visible room differences and clear prices. We treat this as early concept feedback, not a test of the completed prototype.

The newly reviewed transcript is not counted as an additional fourth agent: its relationship to the three-agent summary is unverified.

What the numbers can support

Input	Interpretation and limit
95%	An interview-reported estimate of agent booking within the selected segment; no underlying dataset supplied.
4 of 5	A local guest sample supporting the relevance of assistants; no population estimate.
3 agents	Paraphrased findings in the final deck; the full recruitment and sampling record is incomplete.
11 of 16	A team-coded review sample concerning situational services; neither a hotel complaint rate nor an outcome measure.

Who we could not adequately reach

We did not secure a broad sample of staying guests, the full chain of procurement and travel intermediaries, or property teams across shifts and departments. We had no real post-design user testing with agents or hotel staff. We also lacked operational logs, approved cost data and a verified return-rate baseline.

The supplied GXP material provides an important organisational countercheck: MO already describes preference records, traces, task delegation and service guidance. We therefore position our work as a proposed extension to an existing service system.

Our insight, in four lines

We moved from collecting more preferences to making selected needs easier to express, interpret and fulfil. The four lines below are our evidence-qualified synthesis.

The Assumption

We believed that collecting guests' expressed preferences before arrival would allow the hotel to prepare a room that felt like home from the moment they entered.

The Data

The evidence shows that assistants often mediate bookings in our selected context, while some needs emerge only during the stay. It also shows that trusted calls can already work well, some private needs are deliberately withheld, and MO already has preference records and service guidance.

The Insight

So we now believe that helping these guests feel recognised depends on connecting the needs they choose to share with trusted human judgement and a feasible service response. Optional visual prompts may help assistants express unfamiliar possibilities; Dream Makers may add contextual understanding, provided each action has consent, ownership and delivery limits.

The Invalidation

We would be wrong if existing trusted channels already conveyed the relevant needs adequately, if target assistants preferred those channels to our prototype, or if the added discovery and coordination failed to improve recognition within the hotel's capacity.

How our interpretation changed

Our workshop statement connected guest, agent and staff knowledge around the idea of "feeling at home". The final playback sharpened this into two gaps: needs left unsaid online and situational observations that were difficult to share.

On reviewing the full assistant transcript and GXP deck, we had to qualify that story. More disclosure is not always desirable, and a lack of disclosure is not automatically a failure of texting. We cannot claim that MO has no shared process. Our proposed contribution is a clearer connection between optional expression, situational learning and operational confirmation.

What the service has to connect

Knowledge	The question we carry forward
Assistant's knowledge	What can be shared for this stay, through a channel the guest and assistant trust?
Guest's choice	What is wanted, unwanted, uncertain or deliberately private?
Dream Maker's observation	What is a useful clue, and what still requires clarification?
Hotel's capacity	Who can deliver, by when, at what cost, and who confirms it?

Fig. 2. A reading of the service relationship, redrawn for this case study. These are questions for validation, not confirmed user behaviours.

What we chose, and what we set aside

18 September–early October. We developed two connected touchpoints rather than expanding the preference questionnaire indefinitely.

We chose an assistant-facing visual flow

We used the purpose of the stay to surface potentially relevant room and service choices. A visual room offered a concrete way to discuss arrangements that an assistant might not otherwise think to request. We kept a review step so that a selected option could be checked before submission.

We treated purpose as a starting point rather than a complete description of a guest. The assistant interview challenged our early business-versus-leisure generalisation: familiarity with the guest and the service relationship can matter more than trip category.

We chose a Dream Maker coordination layer

We paired online preparation with in-person discovery. The Dream Maker workspace structures a possible act of care into a plan, approval, departmental feasibility check, execution and review. It is intended to make learning reusable without removing individual judgement.

This adds a proposed workflow around situational care. It does not establish that MO needs its first preference database or its first SOP. The supplied GXP framework already includes profile information, traces and minimum service guidance.

We decided against several directions

We set aside an all-guest, all-property solution. The focused segment and urban context made the design more manageable, but limit how widely our conclusions can travel.

We did not pursue exhaustive mandatory questions or interaction for its own sake. Skipping, saying “not needed” and keeping a trusted human channel remain necessary design conditions.

We did not propose replacing the hotel’s core property system. Any production version would need an agreed connection to existing records and task handling. That integration is unresolved.

We also rejected automatic promises of bespoke service. Suggested plans must remain provisional until the relevant people confirm capacity, cost and timing. A declined or revised plan needs a visible explanation.

Stay Context	Rest	Living	Work	F&B	Bathroom
Business	High	Medium	High	Medium	Low
Honeymoon	High	High	Low	High	High
Anniversary	Medium	High	Low	High	Medium
Family	High	High	Low	High	Medium
Leisure	High	High	Low	Medium	Medium

Fig. 3. Stay-purpose / preference-area matrix in the process file, named 5 October. We used purpose as a starting hypothesis, not a complete description of a guest.

The service as we understood it

We assembled the current journey from interviews, site visits and the supplied GXP “Today” workflow. It is a composite account, not an audited process for every MO property.

Stage	Guest / assistant	Frontstage	Backstage	Where our evidence is limited
Booking	Assistant arranges the stay; other intermediaries may be involved.	Reservations or an established relationship contact receives the brief.	Match profiles, review tier/VIP status, record purpose and preferences.	Our three-role model simplifies corporate procurement, travel platforms and local hosts.
Before arrival	Guest or assistant chooses what to share and how.	Property outreach; Guest Relations / Front Office follows up in a preferred channel.	Review arrivals at 14, 7 and 3 days; delegate traces and service preparation.	No live logs showing response time, handoff loss or the prevalence of rework.
Arrival / rooming	Guest encounters the prepared room and service.	Welcome, check-in and contextual conversation; Dream Maker care where observed.	Coordinate room readiness, amenities and departmental actions.	The Dream Maker observation came from Shenzhen; the role cannot be assumed identical elsewhere.
During the stay	Guest requests, accepts or declines further care.	Staff respond and may notice new preferences.	Update relevant records, follow up and confirm tasks.	We did not observe every department, shift or complete request-to-delivery cycle.

Fig. 4. Composite As-Is journey. The supplied GXP deck distinguishes its existing “Today” process from a proposed “Tomorrow with GXP” process; we keep that distinction.

An existing foundation

MO’s material distinguishes declared, observed and directly stated preferences, reservation notes and actionable traces. It also describes minimum guidance and service tempo. Our earlier “no standard SOP” wording was too broad and is withdrawn.

The narrower opportunity

We could not verify a universal absence of coordination. We instead explored whether selected assistants need a clearer way to express possibilities, and whether situational Dream Maker learning can be linked more explicitly to feasible, owned actions within the existing framework.

 Sources: GXP slides 2, 6, 24–31; interviews; stakeholder map and blueprint in the process file. |

The service we proposed

We proposed a connection from optional expression to an accountable response. A request, an observation, a proposed plan and a confirmed arrangement must remain different states.

Moment	Frontstage experience	Backstage work and proposed owner	Confirmation boundary
1 Express	The assistant explores relevant room options, skips what is unnecessary and reviews the brief.	Reservations links explicit requests to the correct stay and records who supplied them.	Viewing an option does not create a request. Submission is not fulfilment.
2 Clarify	A named contact acknowledges the brief and asks only for missing essentials.	Guest Relations / Front Office checks authority, timing, privacy and existing preferences.	Unknown information stays unknown; no inferred preference becomes a fact.
3 Propose	The Dream Maker notices a possible act of care and clarifies where needed.	Draft the plan; identify resources, cost and timing; obtain appropriate supervisory approval.	An observation is a clue. A plan is provisional until feasibility is confirmed.
4 Coordinate	The contact confirms a deliverable arrangement or explains an alternative.	Departments confirm resources and accept tasks; material changes return for review.	Each action needs an owner, deadline and status. No confirmation while resources are pending.
5 Deliver / learn	The guest receives the service, declines it or requests a change.	Teams record delivery; the Dream Maker captures response and a reusable case, subject to policy.	Record what happened, including rejection and failure; distinguish evidence from interpretation.

Fig. 5. Proposed To-Be service sequence. Ownership and system connections are recommendations; MO has not approved or implemented them through this project.

What changes

The online touchpoint prepares a more legible starting brief. The offline touchpoint supports contextual care and cross-department coordination. We retain a human relationship around both, including acknowledgement, clarification and the option to use an existing channel.

What remains to be worked out

The prototype does not connect to live availability, guest records, budgets or task systems. Permissions, retention, assistant authority, late bookings and room changes still require property-level agreement and operational testing.

Two touchpoints, one service proposition

By early October, our work had become an assistant-facing room flow and a Dream Maker workspace. Both demonstrate interactions; neither proves live service performance.

The assistant's side

We organised choices around the stay's purpose, then room zones and service details. The later mobile flow lets an assistant inspect the room, set relevant preferences and review the stay brief. These choices make possible arrangements visible; they do not replace the conversation with a trusted contact.

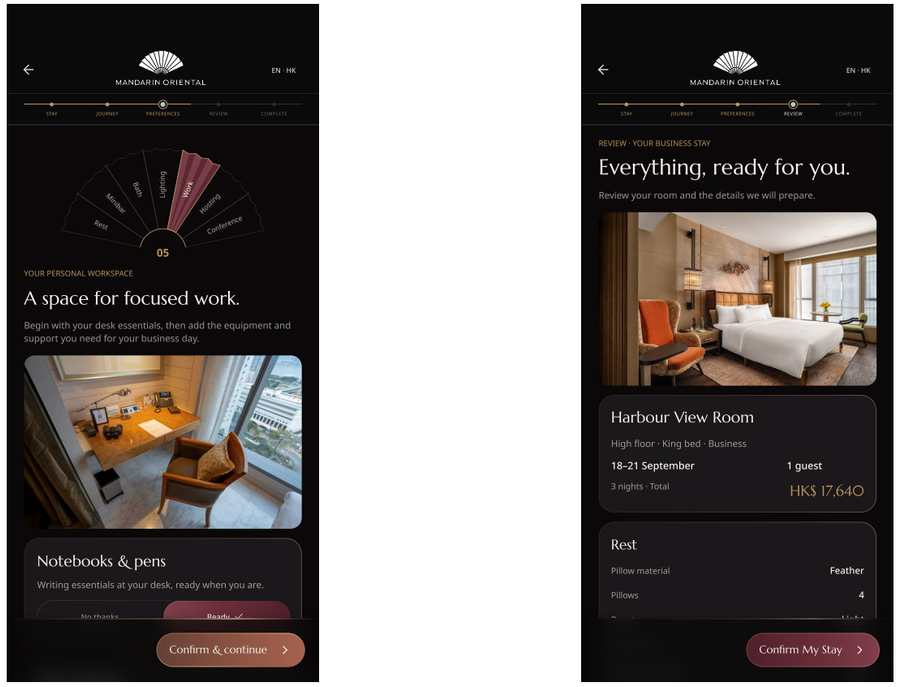

Fig. 6. Work-area preferences and review screens, cloned as editable frames from our final mobile prototype.

Our earlier web prototype explored bed-side selection, temperature and spatial comments. It separated exploration from explicit selection. We still need to test whether the later flow makes skipping, "not needed" and uncertainty clear.

The Dream Maker's side

We structured discovery into a proposed plan, supervisory review, resource confirmation, task coordination and outcome review. Suggestions remain aids to human judgement. The workspace represents a possible operational layer, with mock content and no verified system integration.

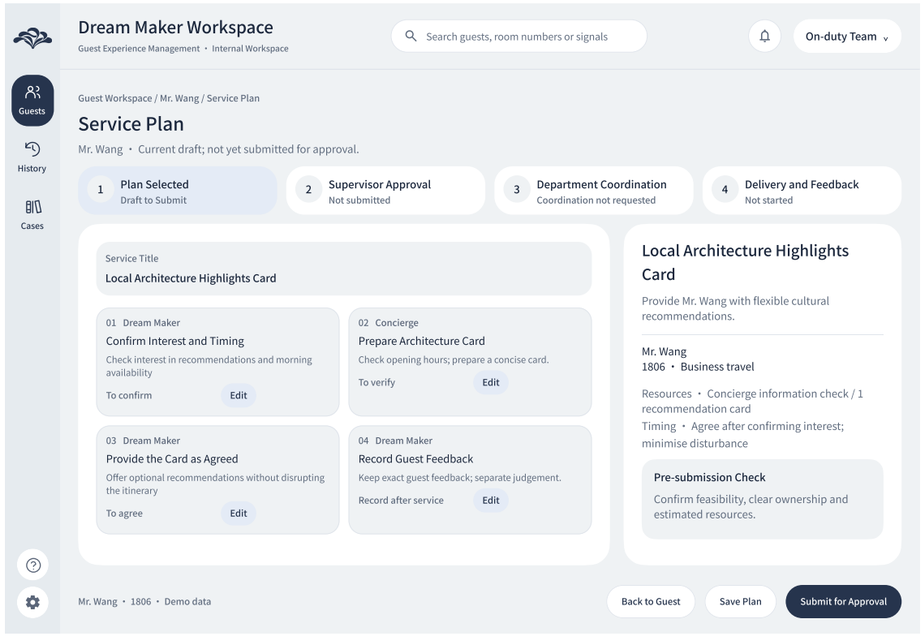

Fig. 7. A service-plan draft in the Dream Maker prototype. It supports planning; it is not evidence that a hotel team accepted or delivered the plan.

We used simple simulated walkthroughs to explore how the interfaces and sequence might work. We did not run real post-design tests with assistants, guests or hotel staff, and did not test the service in a live property. The prototype therefore supports a proposition, not a claim of faster booking, lower workload or increased loyalty.

The numbers and the owner

28 September–2 October. Building the tracker exposed a gap between attractive targets and numbers we could actually defend. We preserve that distinction here.

User-side question	Measure and starting evidence	Business-side companion measure	Proposed owner	Current status
Is expression easier?	Time to communicate an agreed request, plus accuracy. “>5 min” is interview self-report; 3 min was our target.	Total staff minutes per usable, confirmed request, including clarification and rework.	Reservations lead	No measured baseline or result.
Does the guest retain control?	Skip / decline use, misunderstandings and privacy concerns during a task.	Share of accepted requests that are feasible, confirmed and delivered as agreed.	Guest Relations + Reservations	Definitions and baseline needed.
Are handoffs dependable?	Guest or assistant can identify what is confirmed and who to contact.	Task ownership coverage, missed deadlines and unresolved changes.	Front Office + department leads	No live workflow data.
Does care feel appropriate?	Guest response to the delivered care, including unwanted or intrusive service.	Cost to serve and service-recovery incidents; loyalty as a later, separate outcome.	Guest Relations; Finance supports cost	No outcome or cost evidence.

Table 1. A proposed measurement plan reconstructed from the tracker. These role owners have not accepted responsibility through the project.

Claims we cannot carry forward as results

The final deck’s 89.5% → 97.5% capture-rate claim transferred general survey item-nonresponse findings into hotel performance. That inference is unsupported. The review sample’s 11 → 0 complaint target and a 15–25% retention uplift are also untested assumptions, not observed improvements.

How we would count it

We would agree one request definition, start/stop timestamps and a complete workload boundary before a pilot. We would compare the same task through the existing channel and prototype, then track fulfilment separately. Owners must agree baselines, thresholds and reporting cadence before any success claim.

What would have proved us wrong

2 October onward. We wrote failure conditions, but our validation remained a simple simulation. The strongest corrections came from evidence that challenged the framing and the numbers.

What actually broke in our reasoning

Our broad claim that texting prevents needs from being shared did not survive the full assistant account. Trusted telephone coordination could already be efficient, while some information was deliberately withheld. A new interface could add effort or reduce trust instead of helping.

Our “no standard SOP” claim also exceeded the evidence. MO’s supplied material already specifies a preference-to-action framework and minimum guidance. We had to narrow the contribution to a proposed assistant expression flow and a more explicit situational-care coordination layer.

The numerical outcome story broke at the evidence boundary. Survey nonresponse, a small review sample and assumed retention uplift could not establish hotel performance. The case therefore reports these as withdrawn claims or untested targets.

What our simulation did not establish

We did not conduct real post-design user tests or observe live service delivery. We have no reliable task-by-task measurements of time saved, misunderstandings, staff effort or fulfilment failures. Functional interaction checks and a team walkthrough cannot establish desirability or operational feasibility.

The conditions under which we would stop

User side	We would stop or redesign if target assistants preferred the existing channel, refused the interactive flow, or felt it exposed information they did not want recorded.
Business side	We would stop or reduce scope if the flow failed to produce usable, confirmed requests, or raised total staff effort beyond an agreed capacity and cost limit.
The join	We would reject the service claim if more information was collected but arrangements were not delivered as agreed, or the resulting care felt unwanted or failed to improve recognition.

How we would test those conditions next

We would ask real assistants to complete the same request through the existing channel and prototype, then ask property staff to simulate the complete handoff. We would record errors, omitted needs, clarification, total effort and delivery decisions. The hotel and research team must agree failure thresholds before the session.

Our original Class 10 sheet also tied success to return rate. We now treat loyalty as a later hypothesis: an unchanged return rate alone would not identify which part of the service failed.

Landing: what happens next

The course ended with a service proposition and prototypes. The next step is a property-owned validation effort, with a narrow scope and explicit delivery boundaries.

First, agree the service boundary

We would begin with one property, one defined guest group and a small set of deliverable room and service requests. We would map the proposal onto the property's current GXP, profile and trace processes, and identify where it adds value rather than duplicates work.

Before any real guest use, the hotel would need to agree who may act for a guest, what information may be recorded, what requires consent, how long it is retained and who may see it. A human contact must remain available, and guests and assistants must be able to decline the digital route.

Then, validate before extending

We would establish the existing-channel baseline, run comparative assistant tasks and carry out a staff simulation across departments. Late booking, changed room, declined service, unavailable resources and a revised plan should all be included.

Only after those tasks meet agreed criteria would we propose a limited live pilot. We would monitor confirmed-request quality, delivery reliability, total staff effort and guest response. Wider rollout and claims about loyalty or spend would require further evidence.

Who would need to own the work

Role to agree	Proposed responsibility
Property sponsor	Approve scope, staffing capacity, budget and the decision to proceed or stop.
Front Office / Guest Relations	Own the guest-facing promise and the cross-department service journey.
Reservations	Own assistant intake, acknowledgement, clarification and the request-quality baseline.
Department leads	Accept tasks and confirm resources, timing, delivery and exceptions.
IT / data owner	Agree access, integration with existing systems, security and retention.
Finance / service-quality lead	Agree cost-to-serve measures and the balance between service value and operational load.

These are proposed roles, not appointments or commitments obtained during the course.

What we can hand over now

We can hand over the research trail, the reframed brief, the two prototypes, a proposed journey and workflow, the revised tracker and explicit failure conditions. We cannot hand over an approved operating model or demonstrated business outcome. That is the boundary between our completed course work and the work still to be agreed.

The first wall: what we knew and assumed

We began by separating the brief's claims from our own assumptions and the questions that required evidence. The photographs preserve that early uncertainty.

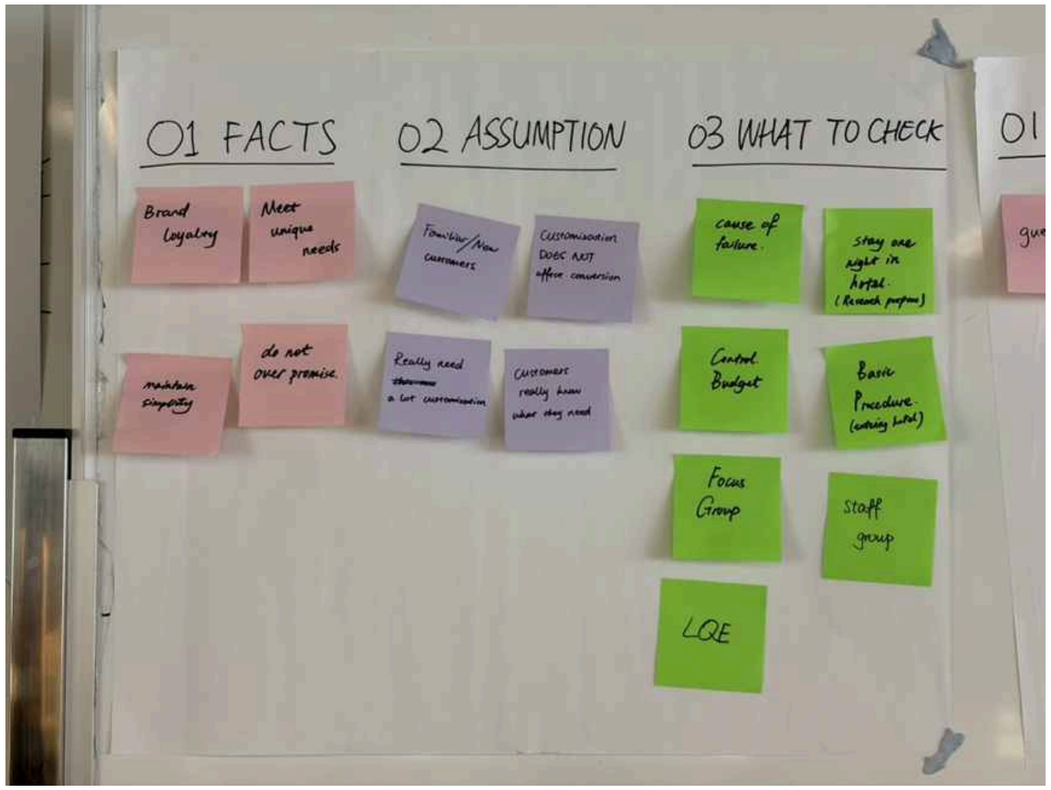

Fig. A1. Facts / Assumptions / What we need to check. Original wall photograph.

Fig. A2. Who is named / who is missing, and our early intent. Original wall photograph.

Why we kept this artefact

The exercise forced us to distinguish what the assigned brief had already stated from what we still needed to learn. Operational complexity was given in the brief. Guest motivations, preferred channels and what "feeling at home" meant were still open questions.

The wall is evidence of our framing process. It is not independent evidence that the assumptions written on its notes were true.

From a guest-only frame to a service relationship

We began with the staying guest, then examined the assistant and the hotel teams who translate and deliver the request. This helped us see why the person using a booking interface might not be the person who experiences the room.

The later assistant interview broadened the chain again: procurement, travel platforms, local hosts and trusted account contacts may also intervene. Our simplified map remains a design model.

Research evidence and its limits

We keep dated source records separate from later synthesis. An interview report, a field observation and a measured service outcome carry different kinds of weight.

Date / source	What the record is	What it contributed	What it does not establish
4 Sep / S02 brief	Assigned Mandarin Oriental brief.	Expectation management, operational complexity, cost to serve and the commercial intent after booking.	An audited NPS model, quantified process baseline or verified current operating performance.
7 Sep / Kate transcript	Proxy stakeholder account.	Context for the halted personalisation effort and the tension between recognition and fulfilment.	A full failure report or direct testimony from all affected departments.
12 Sep / final deck	Five guest conversations reported in the research summary.	Four of five said an assistant arranged their booking; this supported our chosen focus.	A representative share of guests or independent verification of wealth or VIP status.
16 Sep / MO meeting	Further MO interview and meeting summary.	Agent involvement, high-touch service context and operational constraints.	The raw booking dataset behind the reported 95% estimate.
19 Sep / assistant PDF	One machine-transcribed assistant conversation.	Trusted-channel counterexample; privacy concerns; optional menus, visible choices and manager acknowledgement.	Behaviour on the final prototype. Identity linkage to the three-agent summary is not confirmed.
Undated / site & review work	Shenzhen visit, three-agent synthesis and 16-review coding in the final deck.	Dream Maker discovery; selected accounts of coordination and situational service.	Uniform practice across properties, a complaint rate or a service improvement.
Supplied / GXP deck	35-slide internal programme / training material.	Existing preference taxonomy, Today workflow, traces, minimum guidance and a proposed Tomorrow workflow.	Verified rollout status, the stopped pilot’s failure analysis, or live property compliance.
5 Oct / project context	Working migration summary and room-prototype archive.	Evolution of the spatial interface and issues still open at that date.	Final approved requirements, user-test results or live system integration.

We use anonymised research summaries. We omit private third-party details that are irrelevant to the service-design question.

How the research changed our scope

The process did not move directly from one problem to one interface. We repeatedly adjusted who the service was for and what kind of knowledge it needed.

Our early position

We began with the belief that knowing preferences before arrival would help prepare a room that felt like home. That position made preference collection appear to be the central design problem.

The guest and agent research shifted our attention towards the person representing the staying guest. An assistant might know routines and schedules but still need to recognise which hotel arrangements were possible. This gave us a reason to explore visible, concrete choices rather than only an empty text field.

The second shift: care in context

At Shenzhen MO, the Dream Maker role suggested that relevant needs could emerge in a conversation or situation that a pre-arrival form could never fully predict. We widened the service from a collection interface to a relationship between preparation, observation and response.

Our original four-line workshop statement therefore connected guest, agent and staff knowledge and asked what “feeling at home” meant to each guest. The shared guest profile and handoff process became part of the concept, alongside the online touchpoint.

The challenge inside our own evidence

The 19 September assistant transcript prevents a simple “phones and texts are inefficient” conclusion. Familiarity made requests easier; a trusted manager mattered; privacy could make less disclosure the right outcome. We therefore treat visual expression as an option for a specific need, not a universal replacement channel.

The supplied GXP deck also prevents an “MO has no process” conclusion. It documents existing information categories and traces. This case study now describes our proposal as building on that foundation and gives the narrower claims explicit limits.

Working context	The design implication we retained
Before arrival	Help the assistant recognise and communicate selected possibilities; acknowledge the request.
In person	Allow contextual discovery, with clarification and respect for privacy.
Across teams	Translate a chosen response into a feasible plan with an owner and a visible status.
After delivery	Keep outcomes and rejected assumptions available for learning, subject to policy.

This is a retrospective account of the reasoning. It does not imply that every qualification was resolved in the final presentation.

App audit and early room explorations

We examined the existing digital experience and explored a spatial interface. The archive shows interaction possibilities, rather than evidence that these interactions improved a user’s task.

We examined existing preference choices

The existing app capture already shows concrete choices, including a late-checkout time selection. Our work was therefore not the invention of selectable preferences. We explored how room context and stay purpose might help an assistant identify additional relevant arrangements.

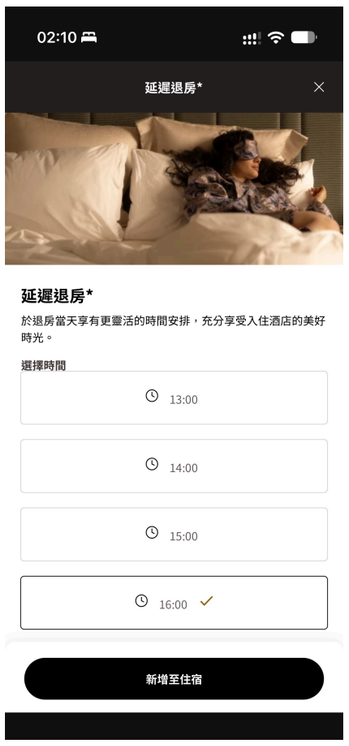

Fig. A3. Existing MO app capture from our audit: late-checkout options. This is research reference material, not our redesigned screen.

The audit alone did not show that guests misunderstood this screen, or that a spatial interface would perform better. That comparison still requires task-based research.

We explored direct interaction with the room

Fig. A4. Early spatial prototype: room hotspots, temperature and preferred bed side.

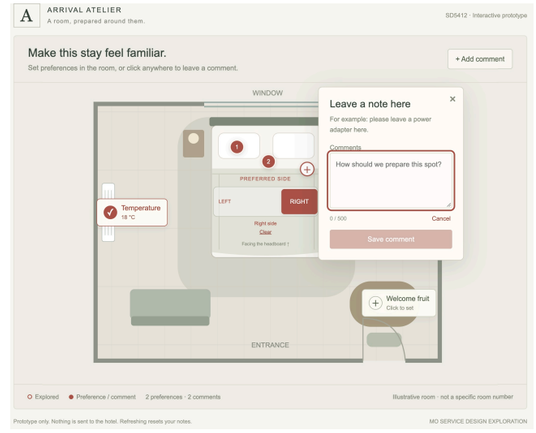

Fig. A5. A later state in the exploration: comments anchored to a location in the room.

We moved editing into the illustration rather than a fixed side panel. We then added bed-side dragging, a temperature wheel and comment editing. Viewed, explicitly selected and not specified were distinct states; exploring a feature did not itself create a service request.

From a stakeholder wall to an information map

We mapped who knows something about the guest, who passes it on and who delivers the service. Drawing a connection did not verify that the handoff worked.

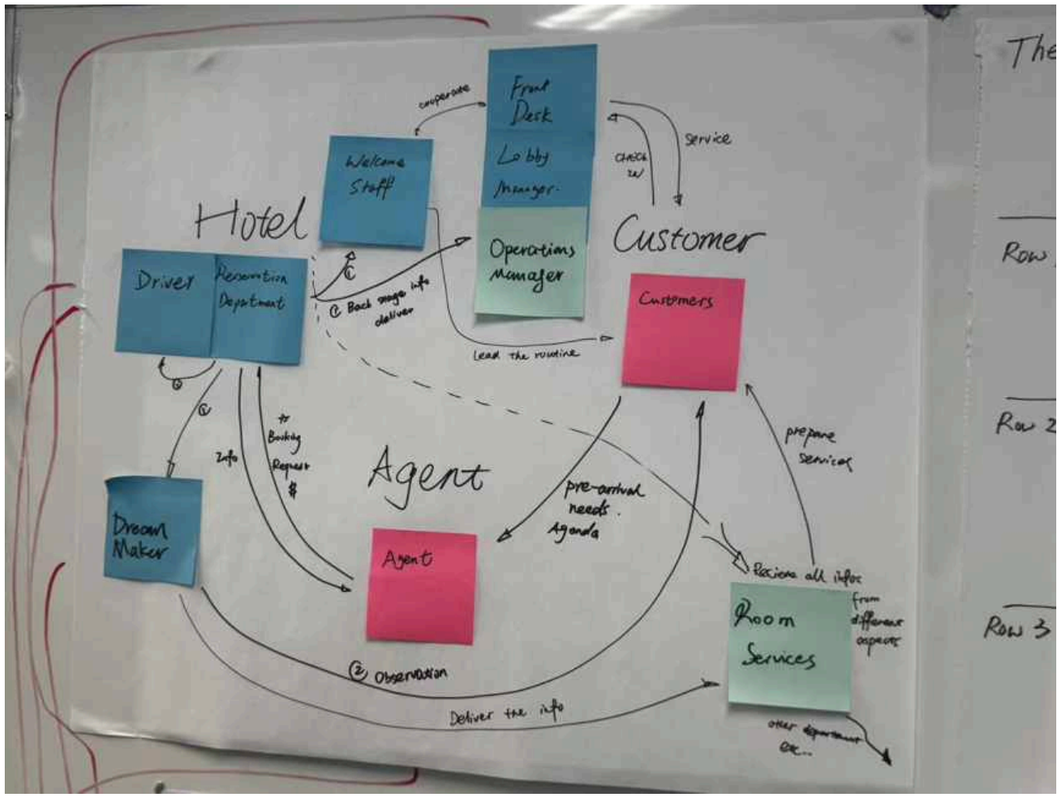

Fig. A6. Original workshop stakeholder map. The photograph preserves the work as it stood.

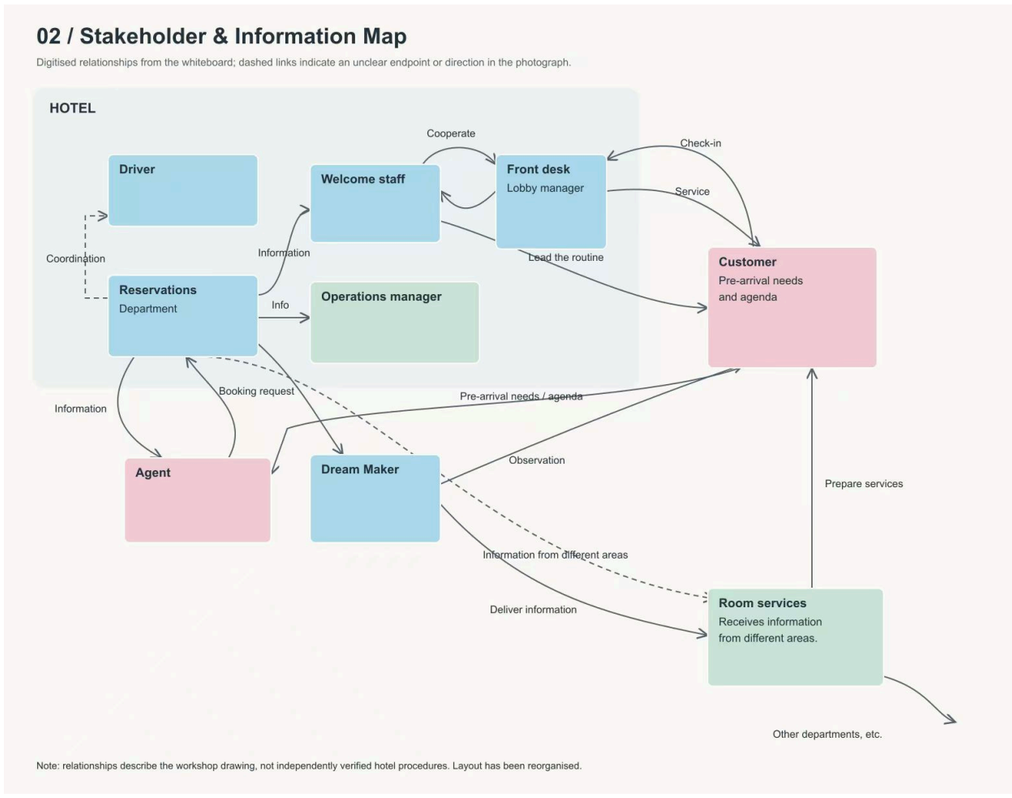

Fig. A7. Existing digitisation of the workshop map. Dashed links preserve uncertainty in the photographed drawing.

What the wall helped us see

The map put the assistant outside the hotel but inside the service relationship. Reservations, welcome staff, the front desk, Dream Maker and service departments could each hold a different part of the guest's context.

This made the handoff itself a design concern. A preference needed both a route into the hotel and a route to the people able to act on it.

What the map still leaves out

The simplified model does not represent every procurement route, travel intermediary, authority boundary or system permission. We did not verify each arrow through real operational records.

The later proposal adds clearer ownership and confirmation states. Those additions are design decisions to be tested, not discoveries already established by this diagram.

The blueprint as a working artefact

We used the journey to connect the visible guest experience with preparation and coordination behind it. The workshop sheet mixed current activity with ideas we wanted to try.

Fig. A8. Original service-blueprint photograph. Handwritten additions show the work developing.

01 / The Journey

Digitised from the whiteboard. English lightly edited for readability; unclear or obscured text is explicitly marked.

PROJECT INTENT — Close the communication gap in pre-arrival service for high-end, wealthy people.

	01 Talk to agent	02 Book hotel	03 Before travel	04 Departure / arrival	05 Check-in / room
	Basic information: travel, personal information, time. Agent develops needs based on their understanding.	Phone / email / app. Customer / boss's needs and membership.	Double-check booking details, if the hotel has responded.	Driver pick-up. Trip departure → arrival.	Check in. Arrive at the room.
ROW 1 Original row label	**Agent / information** Transfer schedule and trip information [end obscured]. Receive requests from boss. Analyse further needs [partly obscured].			**Arrival contact** Picked up by welcome staff. Deal with driver / Dream Maker / butler.	**Room handover** Give new needs to Dream Maker / butler. Go to room. Front desk & Dream Maker [last words obscured].
OVERLAID CONCEPT NOTES — separated here so every visible note can be read	**Playful interaction** "Playing chess to see answers." Small game-grid sketch in the original.	**Open the app** Submit booking details through a playful way.	**Hotel confirmation** Receive confirmation from the hotel.	**On the way to hotel** The boss expresses needs through a playful way.	**Direct to the room** Everything is prepared.
ROW 2 Original row label	**Room preparation** Room service prepares the room. Welcome gifts.	**Record preferences** Record preferences in the system / database.	**Welcome staff / DM** Welcome staff / Dream Maker receive preferences and schedule.	**Prepare services** Schedule / prepare [opening words obscured]. Dining services prepared.	**Further needs** Dream Maker finds further potential needs and delivers information to the system.
			Feedback to system		
ROW 3 Original row label	**Volume / accuracy** Holidays: high arrival volume. Recorded information may not be precise.	**Three pathways** Phone call / email / app.		**Coordination / system** Reservations coordinate profiles across departments. OPERA — "under XXX category" [unspecified].	

Source fidelity: rows retained; overlapping notes reflowed. Cropped purple notes at the bottom cannot be fully transcribed.

Fig. A9. Existing digitisation of the workshop journey. It separates overlaid concept notes and marks obscured text.

From booking to the first room encounter

We followed the guest from discussing travel with an assistant through booking, preparation, pickup, check-in and room entry. The lower rows made preparation, preference records and departmental coordination visible.

This expanded the problem beyond the booking interface: a clear brief is useful only if the hotel can interpret it and deliver the arrangement.

Why the final case separates As-Is and To-Be

The workshop artefact contains proposed playful interactions alongside descriptions of current work. Treating the entire sheet as an observed process would overstate the evidence.

Pages 05–06 therefore distinguish a composite current journey, grounded partly in GXP's Today workflow, from our proposed service. Neither diagram establishes that every MO property operates identically.

Storyboard and the limits of simulation

We used a storyboard to make the proposed service sequence concrete. Its ending is an intended experience, not an observed guest outcome.

Fig. A10. Our six-scene pre-arrival storyboard. Original illustration; undated in the supplied archive.

What we tried to make explicit

Scene	The service hypothesis inside it
Guest → assistant	A conversation establishes travel intent and known routines; the assistant may contribute additional context.
Assistant → interface	Visible options could make possible arrangements easier to recognise and express.
Submission	The assistant can communicate a useful brief without excessive effort.
Hotel confirmation	The hotel distinguishes what can and cannot be arranged before the assistant reports back.
Arrival interaction	A Dream Maker may notice something the pre-arrival brief did not contain. The original scene leaves the method "TBD".
Room encounter	The guest feels appropriately recognised by the preparation and care received. This was the desired ending.

What we can say about the walkthrough

We had only a simple simulated test of the design. We did not recruit real assistants or hotel staff to complete the final service tasks. The archive does not supply a reliable task-level record of timings, misunderstandings or actions that could not be executed.

The storyboard makes assumptions visible: the assistant will engage, a request will reach a decision-maker, arrival contact will be available and the hotel can fulfil the plan. It does not show that those assumptions held.

For a next test we would include refusal, no additional need, unavailable capacity and a late change, then record both the user experience and the complete staff workload. This is a proposed test, not one already conducted.

The assistant prototype evolved through several forms

By 5 October, we had documented a working spatial prototype. The later Figma flow expanded the concept into a mobile stay-planning journey.

Early spatial versions

Version	What changed in the working record
V1	A room plan on the left, fixed editing and summary panel on the right, and simulated submission.
V2	Editing moved into room popovers; comments could be anchored to a chosen location.
V3	English interface; bed-side drag and snap; a vertical temperature wheel; comment add, edit and delete.

Stay Context	Rest	Living	Work	F&B	Bathroom
Business	High	Medium	High	Medium	Low
Honeymoon	High	High	Low	High	High
Anniversary	Medium	High	Low	High	Medium
Family	High	High	Low	High	Medium
Leisure	High	High	Low	Medium	Medium

Fig. A11. Stay-purpose / preference-area matrix in the process file, named 5 October. High / medium / low are our design assumptions.

A distinction we wanted to preserve

A neutral hotspot meant not interacted with; an outlined state meant explored; a filled state meant an explicit preference or comment. The archived source separately stores visited controls and selected values. Functional checks of these states are not evidence of user comprehension.

Later mobile prototype

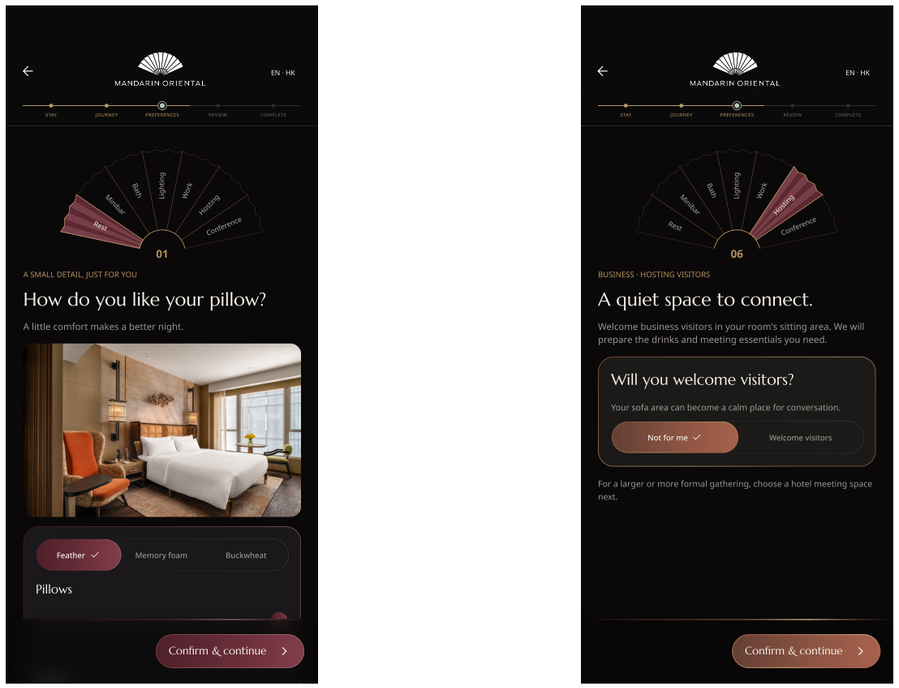

Fig. A12. Rest and in-room meeting screens from the later Business Demo. These are editable copies of the original Figma frames.

We moved from a small room-interaction experiment to a fuller flow covering stay purpose, relevant preference areas, room details and review. The interface demonstrates how information might be structured for an assistant.

We still need to establish whether assistants can distinguish a default from a request, skip without anxiety, communicate uncertainty and understand what requires hotel confirmation. Neither implementation checks nor positive interview reactions to an earlier idea answer those questions.

The Dream Maker workflow became more explicit

We developed the offline idea from noticing and arranging care into a sequence of planning, approval, coordination and review. The prototype makes these states visible.

From an individual idea to a deliverable plan

We proposed that a Dream Maker capture a relevant signal, consider possible responses and draft a plan. Supervisory review and departmental feasibility feedback would then turn a promising idea into an arrangement the hotel could actually deliver.

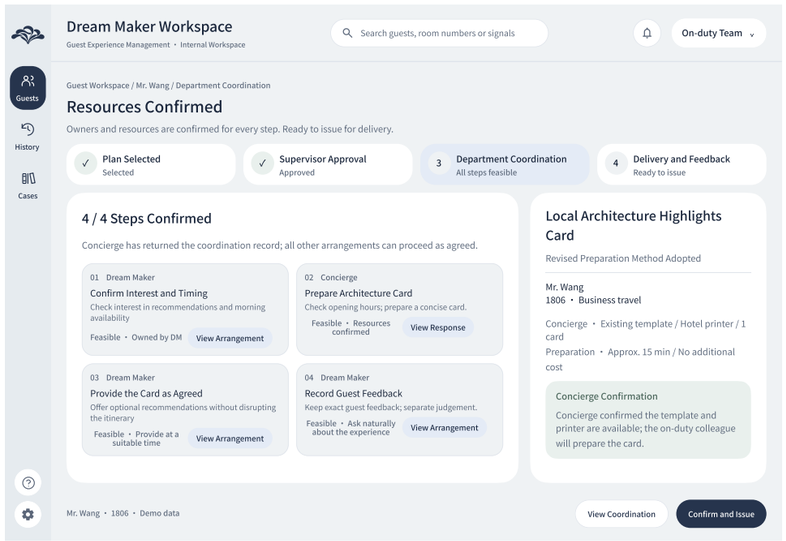

Fig. A13. Resource-confirmation state in our prototype. The content is illustrative, not a record of a real service case.

The workflow also includes delivery, guest response and outcome review. Reusable cases could support future planning, but should preserve what was observed, what was inferred and what the guest actually accepted or declined.

AI-assisted suggestions and case structuring remain proposed support functions. Their appropriateness, data access and effect on staff judgement have not been tested.

Changes must reopen the right decision

We included revision and resubmission states so that a changed plan would not silently inherit an earlier approval. Pending resources, declined departmental support or a material change need a visible route back to review.

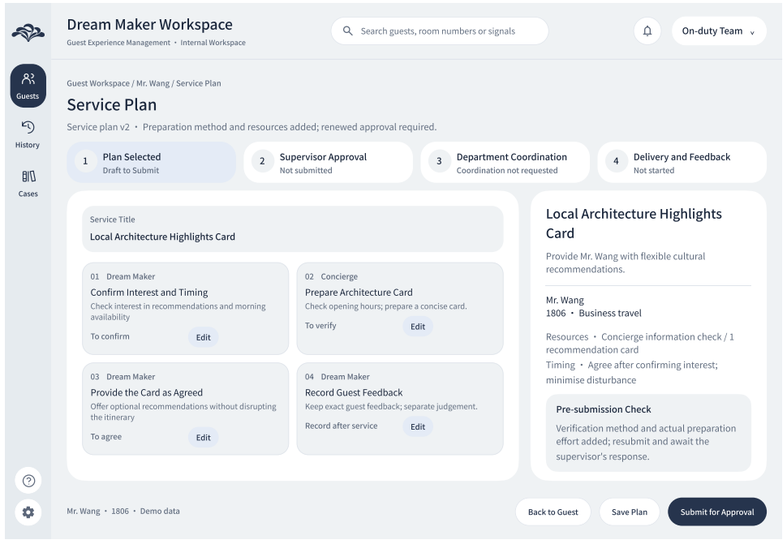

Fig. A14. Modified plan awaiting renewed approval in the prototype. A visible state does not prove that an operational team will follow it.

The implementation question is whether this sequence fits real staffing, shift changes and current task systems. We have not verified that it reduces workload, prevents missed handoffs or improves care. A cross-department staff simulation is the next appropriate test.

This is a proposed extension to the preference and trace framework described in GXP. It should be evaluated against the property's existing process and trusted working practices.

Criteria, source coverage and archive gaps

We retain the unfinished parts of the work because they show where the reasoning still needs evidence. A missing artefact or untested number must remain visible.

The original tracker was a work in progress

Draft item	Status in the case study
10 → 3 min reservation	A Class 10 ambition. The later deck used >5 → 3 min. Neither is a measured before-and-after result.
6 → 1 min confirmation	Draft target without an agreed measurement boundary or observed baseline.
55% → 90% requests	Unverified draft values; more requests alone do not establish more useful or deliverable care.
20% → 55% return rate	Unsupported draft hypothesis; no causal evidence linking this service to that change.
User / business / join criteria	A useful structure, but real task results and agreed failure thresholds are still missing.

What we could not recover as original progress evidence

We did not locate identifiable Crazy 8 sheets or handwritten A5 reflections in the reviewed material. Course instructions establish that these methods were requested, but do not establish which sheets the team completed. We therefore leave the gap explicit.

We also lack a complete task-level simulation log, a verified quantitative baseline, a cost model and property-approved ownership. We do not reconstruct these as if they had existed.

What this account draws on

We reviewed the assigned brief, supplied course material and the project's research, synthesis, pre-presentation, final-presentation, app and Dream Maker pages. We used the original wall photographs, stakeholder map, blueprint and storyboard as progress evidence.

The iCloud review added the 19 September assistant transcript, the 5 October working context, the archived room-prototype source and the 35-slide GXP deck. Its Class materials subfolder was empty at review.

Limits of source coverage

The shared Figma folder index could not be fully enumerated, so we do not claim an exhaustive review of every standalone presentation file in that folder. The known project file and accessible final presentation were reviewed.

The supplied GXP deck is programme / training material with a Today and a Tomorrow workflow. Its header date and programme title do not establish when a particular property implemented it. We did not obtain a clearly identified failure-analysis deck for the stopped hyper-personalisation pilot.

Reading this appendix

Original photographs and illustrations preserve historical material. Digitisations are labelled as such. New narrative, journey tables and measurement proposals are retrospective synthesis. None should be mistaken for real post-design user-test evidence or approved MO operations.